

1 Introduction

1.1 Distributed Systems Definition

A distributed system in its simplest definition is a group of computers working together as to appear as a single computer to the user.

1.2 Why Distributed Systems

- Scaling
  - Vertical: more memory, faster CPU
  - Horizontal: more machines
- Economics
  - Initially scaling vertically is cheaper until max HW
  - Current x86 max: 64 cores
- Location
  - Everything gets faster, latency stays
  - Physically bounded by the speed of light
  - New Protocols can decrease RT
  - Place services closer to user
- Fault tolerance
  - Every hardware will crash eventually

1.3 Scaling

1.3.1 Horizontal

- Pros
- Lower cost with massive scale
  - Easier to add fault-tolerance
  - Higher availability

- Cons
- Adaption of software required
  - More complex system, more components involved

1.3.2 Vertical

**Moore’s Law:** Nr. of transistors doubles every 2 years.  
**Nielsen’s Law:** High-end user’s connection speed grows by 50% per year.  
**Kryder’s Law:** Disk density doubling every 13 month.  
*Bandwidth grows slower than computer power*

- Pros
- Lower cost with small scale
  - No adaption of software required
  - Less administrative effort

- Cons
- HW limits for scaling
  - Risk of HW failure causing outage
  - More difficult to add fault tolerance

1.4 Distributed Systems Categorization

- Tightly Coupled**
- Processing Elements have access to a common memory

- Loosely Coupled (this lecture)**
- Processing Elements have NO access to a common memory

- Homogenous System**
- All processors are of the same type

- Heterogeneous System (this lecture)**
- Processors of different types

- Small Scale**
- WebApp + database

- Large Scale (this lecture)**
- More than 2 machines

- Decentralized**
- Distributed in the technical sense, but not owned by one actor

1.4.1 Controlled Distributed Systems

- 1 responsible organization
- Low churn
- Secure environment
- High availability
- Homogenous / Heterogeneous
- Examples: Amazon DynamoDB, Client/Server

- Mechanisms that work well:**
- Consistent hashing
  - Master nodes, central coordinator

- Network is under control or client/server**
- no NAT issues

- Consistency**
- Leader election (Zookeeper, Paxos, Raft)

- Replication principles**
- More replicas: higher availability / reliability / performance / scalability
  - Requires maintaining consistency in replicas

**Transparency principles apply**

1.4.2 Fully Decentralized Systems

- N responsible organizations
- High churn
- Hostile environment
- Unpredictable availability
- Heterogeneous
- Examples: BitTorrent, Blockchain

- Mechanisms that work well:**
- Consistent hashing (DHTs)
  - Flooding/broadcasting - Bitcoin

**NAT and direct connectivity huge problem**

- Consistency**
- Weak consistency: DHTs
  - Proof of work

**Replication / Transparency principles apply**

1.4.3 CAP theorem

A distributed data store cannot simultaneously be consistent, available and partition tolerant.

- **Consistency:** Every node has the same consistent state
- **Availability:** Every non-failing node always returns a response
- **Partition Tolerant:** The system continues to be consistent even when network partitions

- Examples:**
- Network partition: AP or CP
  - Blockchain: CP or AP
  - Cassandra (Apple): AP, can be configured CP

1.5 Transparency in DS

**Distributed system should hide its distributed nature**

- Location: users should not be aware of the physical location
- Access: users should access resources in a single, uniform way
- Migration, relocation: users should not be aware, that resources have moved
- Replication: Users should not be aware about replicas, it should appear as a single resource
- Concurrency: users should not be aware of other users
- Failure: Users should be aware of recovery mechanisms
- Security: Users should be minimally aware of security mechanisms